

Tips de monitoreo con SQL



POSTGRESQL – MYSQL

INSTRUCTOR: EMANUEL CALVO

Premisas



CONCEPTOS BÁSICOS

Premisas



- Evitar/monitorear la latencia de escritura y lectura.
 - ✦ Es lo que tardan los dispositivos de almacenamiento, en devolver los datos.
- Aumentar/monitorear el rendimiento de salida de resultados.
 - ✦ Capacidad de respuesta del servidor.
- Seguridad.

*Bases de datos veloces,
implican aplicaciones
Veloces.*

Como lograr ...?



- Menor latencia de I/O:
 - ✦ Más discos.
 - ✦ Discos más rápidos.
 - ✦ Tablespaces y particiones separadas entre si.
 - ✦ Utilización de sistemas de ficheros más avanzados.
- Rendimiento:
 - ✦ Más memoria.
 - ✦ Más y mejores procesadores y núcleos.
 - ✦ Mejor conexión entre servidores en red.
 - ✦ Servidores dedicados.

Monitoreo Básico



- Desde el Sistema Operativo.
 - ✦ Se obtiene el rendimiento en términos específicos del sistema operativo y hardware.
- Desde el motor de Base de datos.
 - ✦ Se puede obtener datos de accesos a objetos y cantidad de datos en caché de los mismos.

Monitoreo de Base



- Estado de los accesos a relaciones.
- Estado de las estadísticas.
- Estados de los índices.
- Estado del caché.
- Procesos en ejecución.
- TPS con pgbench, medir rendimiento.
- Consultas lentas.

Tamaño de la base y tablas



Postgresql

- Select
`pg_size_pretty(pg_database_size(name));`
- SELECT
`pg_size_pretty(pg_total_relation_size(tabla));`

Mysql

- SELECT table_schema "Data Base Name", sum(data_length + index_length) / 1024 / 1024 "Data Base Size in MB"
FROM information_schema.TABLES
GROUP BY table_schema ;
- SELECT table_schema "Data Base Name",
sum(data_length + index_length) / 1024 / 1024 "Data Base Size in MB",
sum(data_free) / 1024 / 1024 "Free Space in MB"
FROM
information_schema.TABLES
GROUP BY table_schema ;

Tamaño tablas



Postgresql

- **SELECT**
pg_size_pretty(pg_total_relation_size(tabla));
- **SELECT**
pg_size_pretty(pg_relation_size(tabla));

Mysql

- **SELECT** table_name,
table_rows, data_length,
index_length,
round(((data_length +
index_length) / 1024 /
1024),2) "Size in MB"
FROM
information_schema.TA
BLES WHERE
table_schema =
"schema_name";

Procesos



Postgresql

- `SELECT * FROM pg_stat_activity;`

Mysql

- `SHOW PROCESSLIST;`
- `SHOW STATUS LIKE '%threads%';`
- `show session status like 'connections';`

Mysql



APARTADO DE CONSULTAS PARA MYSQL

Engines (Mysql)



- Show engines;
- show engine innodb status;

InnoDB Status

Per second averages calculated from the last 59 seconds

-----BACKGROUND THREAD-----

srv_master_thread loops: 4 1_second, 0 sleeps, 0 10_second, 5 background, 5 flush

srv_master_thread log flush and writes: 0

-----SEMAPHORES-----

OS WAIT ARRAY INFO: reservation count 3, signal count 3

Mutex spin waits 1, rounds 30, OS waits 0

RW-shared spins 3, rounds 90, OS waits 3

RW-excl spins 0, rounds 0, OS waits 0

Spin rounds per wait: 30.00 mutex, 30.00 RW-shared, 0.00 RW-excl

-----TRANSACTIONS-----

Trx id counter 502

Purge done for trx's n:o < 32C undo n:o < 0

History list length 7

LIST OF TRANSACTIONS FOR EACH SESSION:

---TRANSACTION 0, not started, OS thread id 2988

MySQL thread id 6, query id 110 localhost 127.0.0.1 root

show engine innodb status

---TRANSACTION 501, not started, OS thread id 3568

MySQL thread id 4, query id 105 localhost 127.0.0.1 root

---TRANSACTION 0, not started, OS thread id 3552

MySQL thread id 2, query id 59 localhost 127.0.0.1 root

-----FILE I/O-----

I/O thread 0 state: wait Windows aio (insert buffer thread)

I/O thread 1 state: wait Windows aio (log thread)

I/O thread 2 state: wait Windows aio (read thread)

I/O thread 3 state: wait Windows aio (read thread)

I/O thread 4 state: wait Windows aio (read thread)

I/O thread 5 state: wait Windows aio (read thread)

I/O thread 6 state: wait Windows aio (write thread)

I/O thread 7 state: wait Windows aio (write thread)

I/O thread 8 state: wait Windows aio (write thread)

I/O thread 9 state: wait Windows aio (write thread)

Pending normal aio reads: 0 [0, 0, 0, 0], aio writes: 0 [0, 0, 0, 0],

ibuf aio reads: 0, log i/o's: 0, sync i/o's: 0

Pending flushes (fsync) log: 0; buffer pool: 0

191 OS file reads, 7 OS file writes, 7 OS fsyncs

0.00 reads/s, 0 avg bytes/read, 0.00 writes/s, 0.00 fsyncs/s

-----INSERT BUFFER AND ADAPTIVE HASH INDEX-----

Ibuf: size 1, free list len 0, seg size 2, 0 merges

merged operations:

insert 0, delete mark 0, delete 0

discarded operations:

insert 0, delete mark 0, delete 0

Hash table size 195193, node heap has 0 buffer(s)

0.00 hash searches/s, 0.00 non-hash searches/s

---LOG---

Log sequence number 3308699

Log flushed up to 3308699

Last checkpoint at 3308699

0 pending log writes, 0 pending chkp writes

10 log i/o's done, 0.00 log i/o's/second

-----BUFFER POOL AND MEMORY-----

Total memory allocated 49971200; in additional pool allocated 0

Dictionary memory allocated 18671

Buffer pool size 3008

Free buffers 2829

Database pages 179

Old database pages 0

Modified db pages 0

Pending reads 0

Pending writes: LRU 0, flush list 0, single page 0

Pages made young 0, not young 0

0.00 youngs/s, 0.00 non-youngs/s

Pages read 179, created 0, written 1

0.00 reads/s, 0.00 creates/s, 0.00 writes/s

No buffer pool page gets since the last printout

Pages read ahead 0.00/s, evicted without access 0.00/s

LRU len: 179, unzip_LRU len: 0

I/O sum[o]:cur[o], unzip sum[o]:cur[o]

-----ROW OPERATION-----

0 queries inside InnoDB, 0 queries in queue

1 read views open inside InnoDB

Main thread id 796, state: waiting for server activity

Number of rows inserted 0, updated 0, deleted 0, read 11

0.00 inserts/s, 0.00 updates/s, 0.00 deletes/s, 0.00 reads/s

Explorando tablas



```
mysql> select * from  
information_schema.tables  
where TABLE_NAME like 'prueba' limit  
1\G
```

***** 1. row *****

```
TABLE_CATALOG: def  
TABLE_SCHEMA: mysql  
TABLE_NAME: prueba  
TABLE_TYPE: BASE TABLE  
ENGINE: InnoDB  
VERSION: 10  
ROW_FORMAT: Compact  
TABLE_ROWS: 5  
AVG_ROW_LENGTH: 3276  
DATA_LENGTH: 16384  
MAX_DATA_LENGTH: 0
```

```
INDEX_LENGTH: 0  
DATA_FREE: 0  
AUTO_INCREMENT: 6  
CREATE_TIME: 2010-11-  
03 12:09:45  
UPDATE_TIME: NULL  
CHECK_TIME: NULL  
TABLE_COLLATION:  
latin1_swedish_ci  
CHECKSUM: NULL  
CREATE_OPTIONS:  
TABLE_COMMENT:  
1 row in set (0.00 sec)
```

Explorando Tablas (2)

- show table status from mysql like 'prueba'\G

Particiones



```
mysql> select * from
information_schame.partitions where
table_name = 'prueba' limit 1\G
***** 1. row *****
TABLE_CATALOG: def
TABLE_SCHEMA: mysql
TABLE_NAME: prueba
PARTITION_NAME: NULL
SUBPARTITION_NAME: NULL
PARTITION_ORDINAL_POSITION: NULL
SUBPARTITION_ORDINAL_POSITION: NULL
PARTITION_METHOD: NULL
SUBPARTITION_METHOD: NULL
PARTITION_EXPRESSION: NULL
SUBPARTITION_EXPRESSION: NULL
PARTITION_DESCRIPTION: NULL
TABLE_ROWS: 5
AVG_ROW_LENGTH: 3276
DATA_LENGTH: 16384
```

```
MAX_DATA_LENGTH: NULL
INDEX_LENGTH: 0
DATA_FREE: 0
CREATE_TIME: 2010-11-03 12:09:45
UPDATE_TIME: NULL
CHECK_TIME: NULL
CHECKSUM: NULL
PARTITION_COMMENT:
NODEGROUP:
TABLESPACE_NAME: NULL
1 row in set (0.02 sec)
```

Estadísticas



```
mysql> select * from statistics
where table_name like 'prueba' limit 1\G
***** 1. row *****
TABLE_CATALOG: def
TABLE_SCHEMA: mysql
  TABLE_NAME: prueba
  NON_UNIQUE: 0
INDEX_SCHEMA: mysql
  INDEX_NAME: PRIMARY
SEQ_IN_INDEX: 1
COLUMN_NAME: the_key
  COLLATION: A
CARDINALITY: 5
  SUB_PART: NULL
  PACKED: NULL
  NULLABLE:
INDEX_TYPE: BTREE
  COMMENT:
INDEX_COMMENT:
1 row in set (0.02 sec)
```


Working example

```
mysql> show table status like 'prueba'\G
```

Name: prueba

Engine: InnoDB

Version: 10

Row_format: Compact

Rows: **9703**

Avg_row_length: 35

Data_length: **344064**

Max_data_length: 0

Index_length: 0

Data_free: 0

Auto_increment: 16371

Create_time: 2010-11-03 12:09:45

Update_time: NULL

Check_time: NULL

Collation: latin1_swedish_ci

Checksum: NULL

```
mysql> delete from prueba where a between  
80 and 81; Query OK, 218 rows affected (0.30  
sec)
```

```
mysql> show table status like 'prueba'\G
```

Name: prueba

Engine: InnoDB

Version: 10

Row_format: Compact

Rows: **10823**

Avg_row_length: 31

Data_length: **344064**

Max_data_length: 0

Index_length: 0

Data_free: 0

Auto_increment: 16371

Create_time: 2010-11-03 12:09:45

Update_time: NULL

Check_time: NULL

Collation: latin1_swedish_ci

Checksum: NULL

Working example (2)

```
mysql> optimize table prueba\G
```

```
1. row *****
```

```
Table: mysql.prueba
```

```
Op: optimize
```

```
Msg_type: note
```

```
Msg_text: Table does not support  
optimize, doing recreate + analyze  
instead
```

```
2. row *****
```

```
Table: mysql.prueba
```

```
Op: optimize
```

```
Msg_type: status
```

```
Msg_text: OK
```

```
2 rows in set (2.13 sec)
```

```
mysql> show table status like 'prueba'\G
```

```
1. row *****
```

```
Name: prueba
```

```
Engine: InnoDB
```

```
Version: 10
```

```
Row_format: Compact
```

```
Rows: 10508
```

```
Avg_row_length: 31
```

```
Data_length: 327680
```

```
Max_data_length: 0
```

```
Index_length: 0
```

```
Data_free: 0
```

```
Auto_increment: 13300
```

```
Create_time: 2010-11-03 12:09:45
```

```
Update_time: NULL
```

```
Check_time: NULL
```

```
Collation: latin1_swedish_ci
```

```
Checksum: NULL
```

```
Create_options:
```

Postgresql



APARTADO CONSULTAS PARA POSTGRESQL

Monitoreo de Base



- Estado de los accesos a relaciones.
 - Pg_statio_user_tables
 - Pg_stat_user_tables
 - Tamaños
 - ✦ `select pg_size_pretty(pg_database_size('ejemplo'));`
 - ✦ `select pg_size_pretty(pg_relation_size('datos'::regclass));`

Monitoreo de Base



- Estado de las estadísticas.
 - Pg_stats
 - `SELECT * FROM pg_stats WHERE tablename = 'tabla' AND attname = 'columna';`

Monitoreo de Base



- Estados de los índices.
 - Pg_stat_user_indexes
 - Pg_statio_user_indexes

Monitoreo de Base



- Estado del caché.
 - Cantidad de bloques leídos:
 - ✦ `Select pg_stat_get_db_blocks_fetched((select datid from pg_stat_database where datname = 'pampabs'));`
 - Bloques leídos y en caché:
 - ✦ `Select pg_stat_get_db_blocks_hit((select datid from pg_stat_database where datname = 'pampabs'));`

Consultas lentas



Postgresql

- Consultas lentas.
 - Activar `log_min_duration_statement` (milisegundos).

Mysql

- Consultas lentas
 - `log_slow_queries`
 - `log_queries_not_using_indexes`

Trucos



- Realizando un ALTER TABLE sin modificar la tabla se obliga la reestructuración.
- Ordenar los registros físicamente:
 - ✦ `SELECT * INTO tabla2 FROM tabla ORDER BY columna;`

Contribs Postgresql



ALGUNOS CONTRIBS PARA MONITOREO

pg_stat_statements



- Loguea absolutamente todas las consultas.
- Funciones:
 - `pg_stat_statements_reset()`
- Se utiliza:
 - `Select * from pg_stat_statements;`

#Configuracion en el Postgresql.conf:

```
shared_preload_libraries='pg_stat_statements'
```

```
custom_variable_classes='pg_stat_statements'
```

```
pg_stat_statements.max = 10000
```

```
pg_stat_statements.track = all
```

```
pg_stat_statements.save = on
```

pgfouine



- Descargar desde *pgfoundry.org*.
- Necesita un php-cli.
- Si tenemos que usar stderr, el log_line_prefix debe establecerse en:
 - '%t [%p]: [%l-1]'

```
Php pgfouine.php -format html-  
with-graphics -logtype stderr -  
file      <archivo_log>      >  
result.html
```

pgstattuple



- Varias funciones de medición estadística.
- Solo hacen un bloqueo de lectura.
- `pgstattuple('table')`
- `pgstatindex('index')`
- `pg_relpages`

pgrowlocks



- Devuelve la información por tupla.
- Devuelve lock_type.

```
SELECT *  
FROM  
    general g JOIN  
    pgrowlocks('general') p  
    ON (g.ctid = p.locked_row);
```



Esta presentación está
bajo licencia GPL.

Contacte a: postgres.arg
(at) gmail.com

Skype: emanuel.cfranco

Gracias Por asistir