



PALOMINO DB

Proven Database Excellence

PostgreSQL 9.1
FOSDEM 2012

Presentado por:
Emanuel Calvo @3manuek
DBA

www.palominodb.com



About the author and the talk:

- Operational DBA at PalominoDB
- Consultor Principal en Helpame
- DBA Mysql y Oracle en Aedgency.
- Twitter: @3manuek
- Mail: emmanuel@palominodb.com
- The main source of this presentation is the following article:
 - <http://palominodb.com/blog/2011/09/28/postgresql-91-part-1-general-features>




PALOMINO^{DB}

Proven Database Excellence

New features

- Synchronous replication
- Foreign Data Wrappers (AKA SQL/MED)
- Collation per column
- Serializable Snapshot Isolation
- Unlogged tables
- Writable Common table Expressions
- KNN for GiST and GIN
- SE-LINUX support
- Upgrade Python version server side
- PGXN
- More stuff...



PALOMINO^{DB}

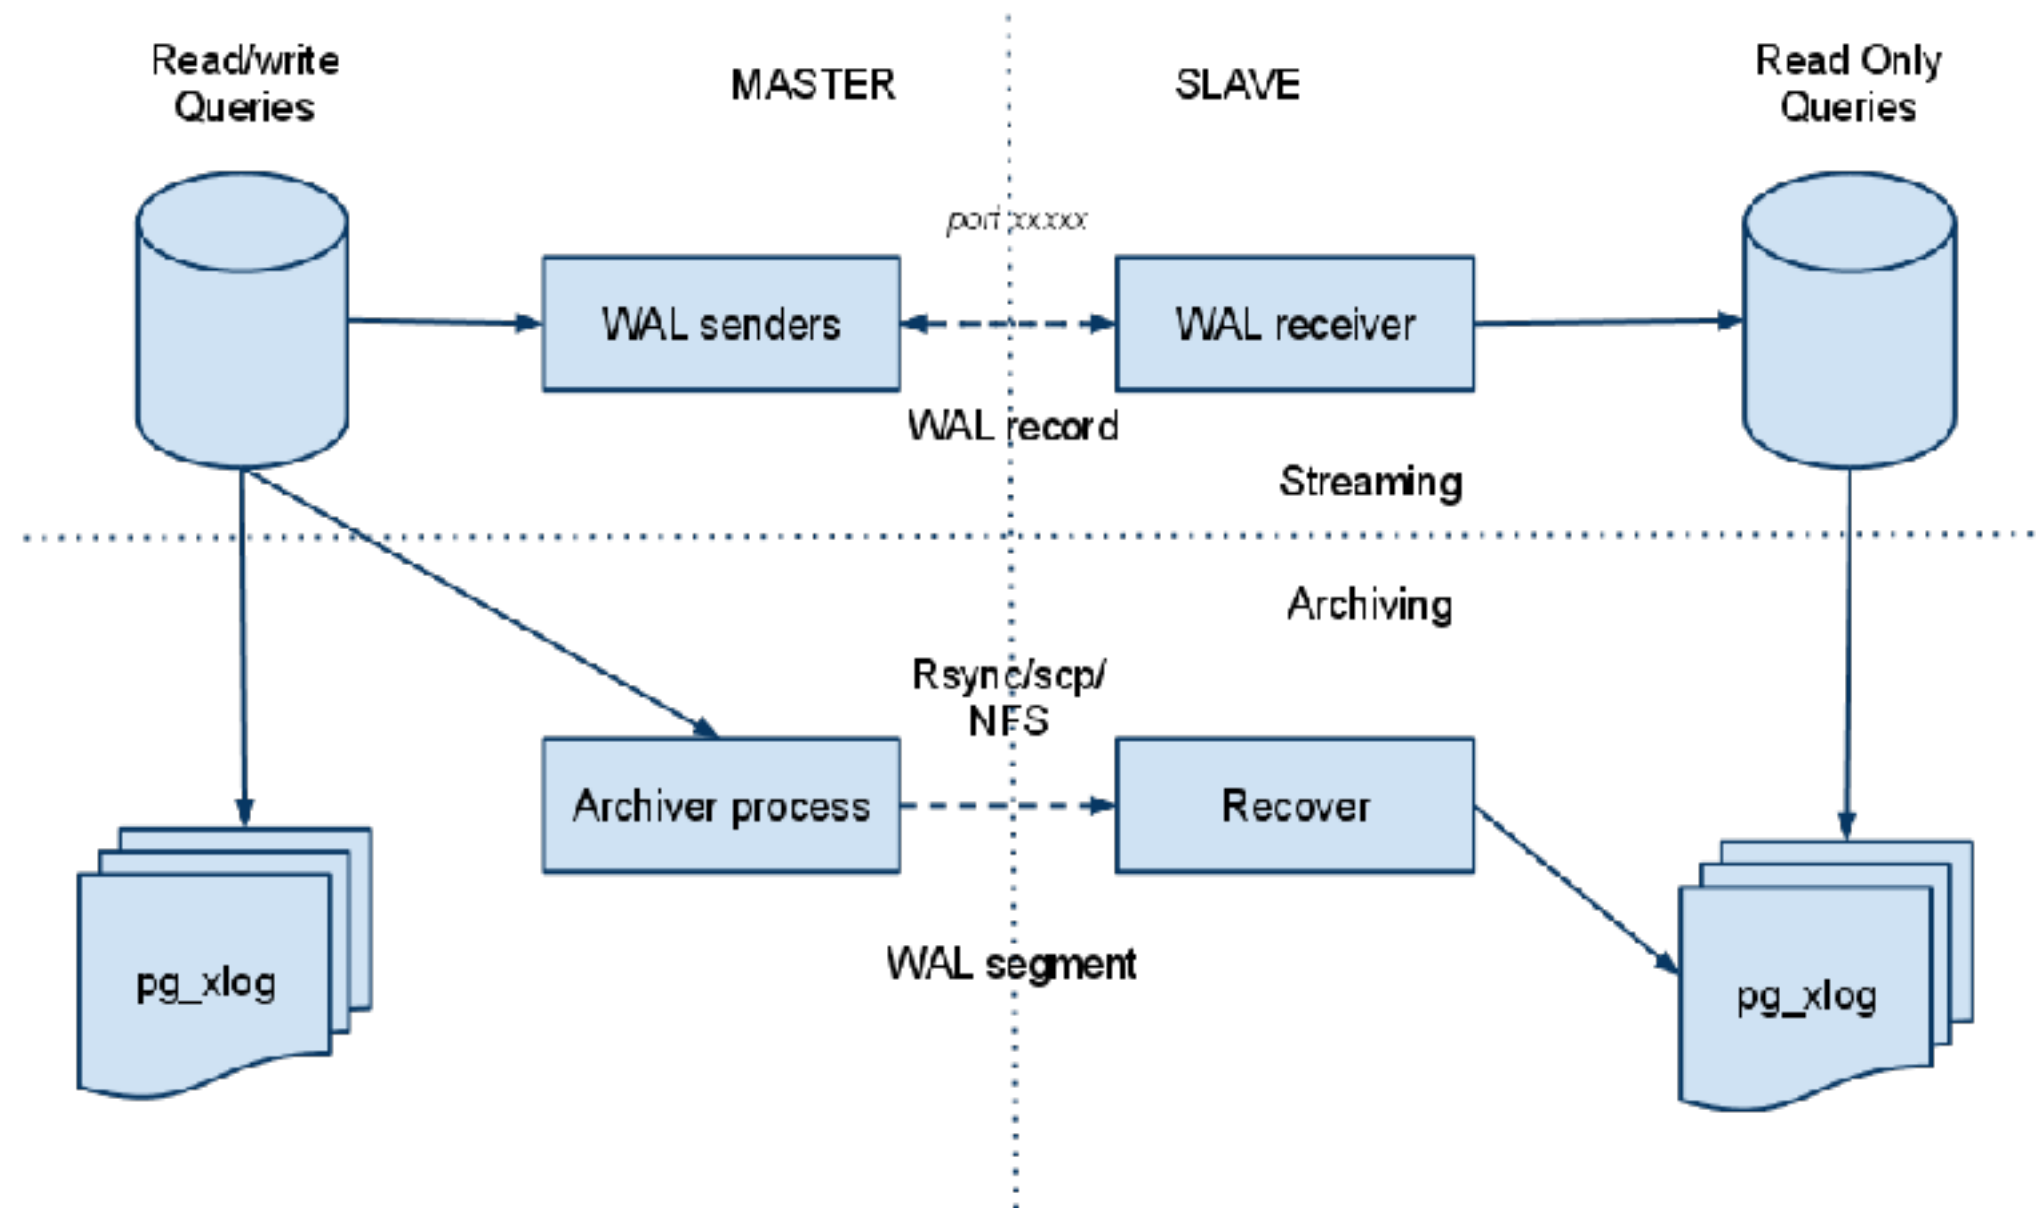
Proven Database Excellence

Synchronous Replication

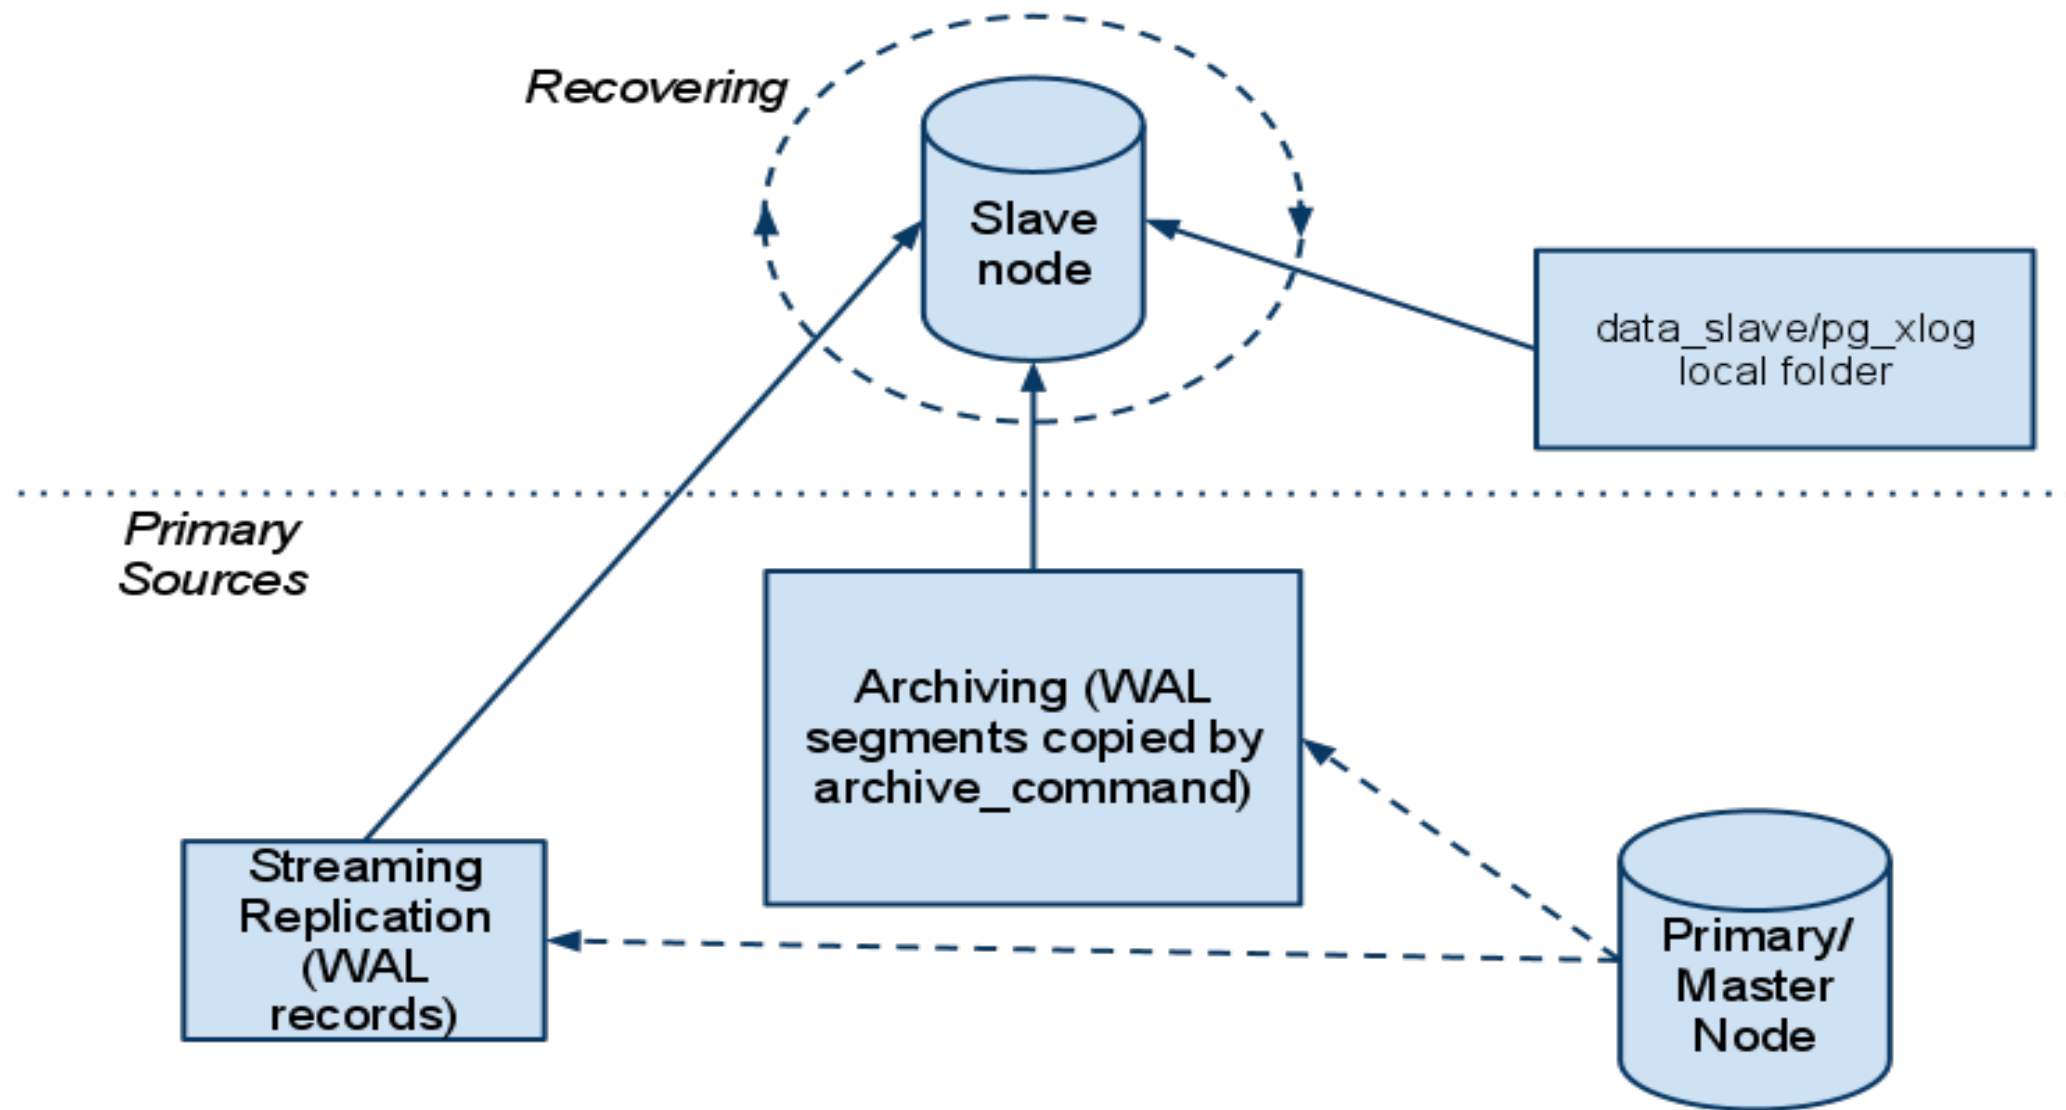
- Control synchronicity per transaction switching `synchronous_commit`.
- Improvements on the 9.0 replication.
- Administrative functions to resume, pause and show state of the slave
 - `pg_xlog_replay_pause()`
 - `pg_is_xlog_replay_paused()`
 - `pg_xlog_replay_resume()`
 - `pg_last_xact_replay_timestamp()`
 - `Pg_last_xact_replay_location()`
- `pg_ctl` now supports “promote” for the slave
- Clone your slave without downtimes or locking (`pg_basebackup`)
- Information about slaves into the master (`pg_stat_replication`).
- [pg_stat_database_conflicts](#) Shows information related queries and vacuum conflicts between master and slave.



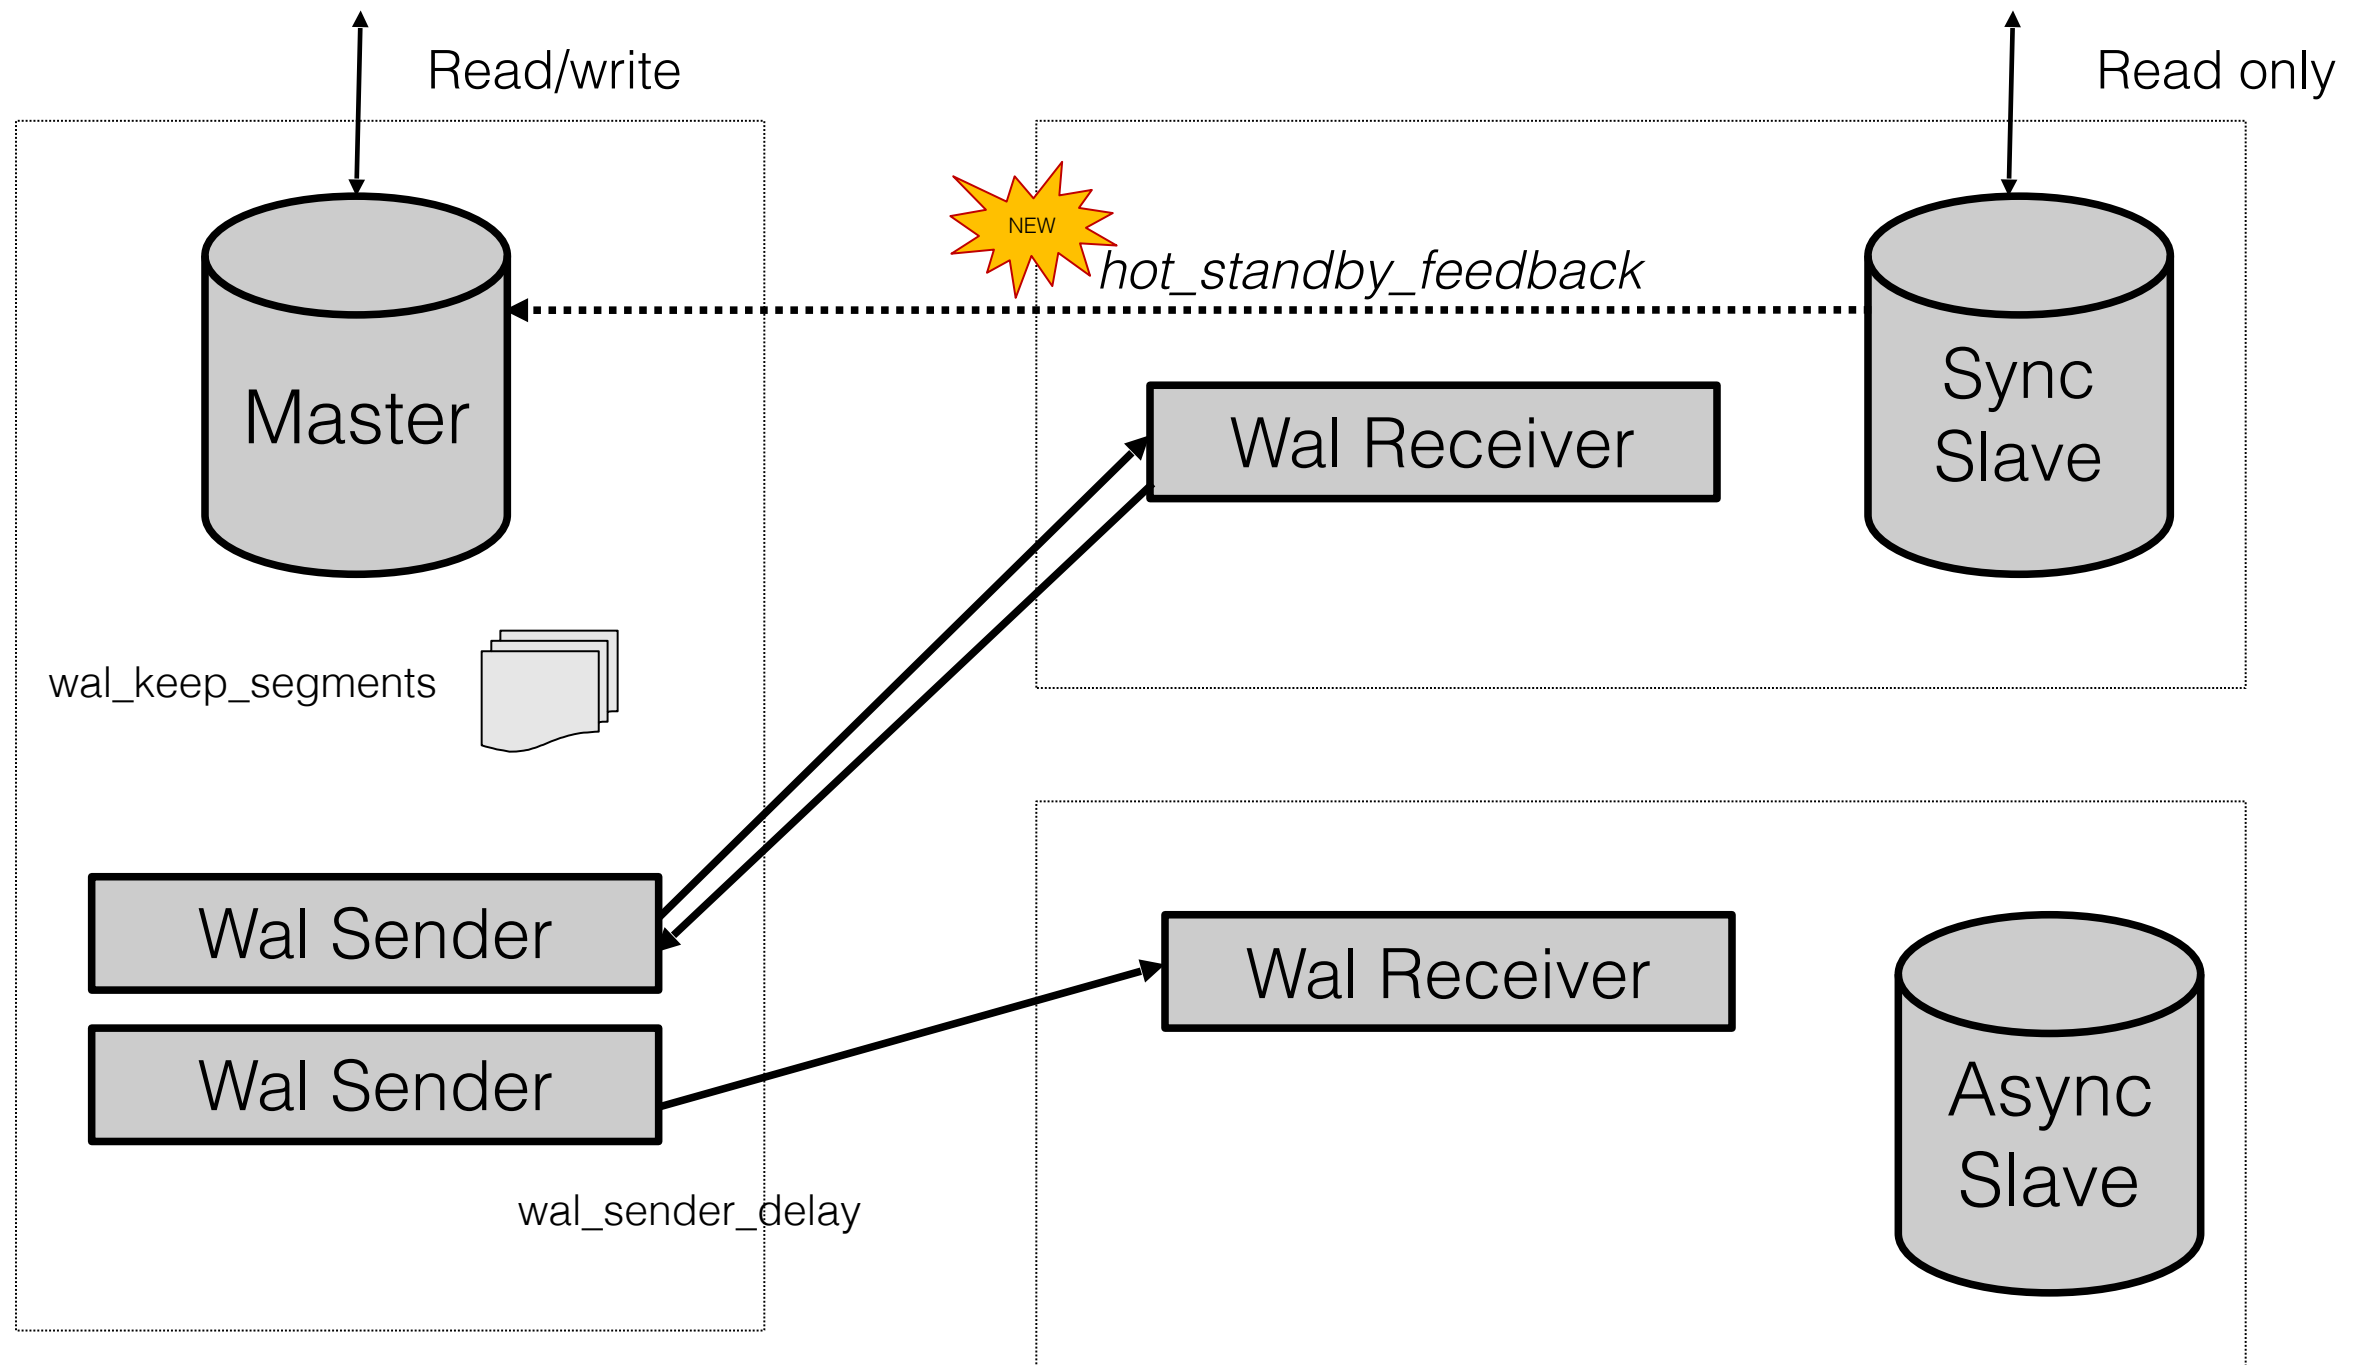
Replication basics



Standby sources



Synchronous replication (deeper scope)



Replication limits

- Only one server can be synchronous.
- Don't support multi master.
- Don't support partial replication.
- Don't support distributed data.



Foreign Data Wrapper (MED – Management of External Data)

```
CREATE EXTENSION file_fdw;
```

```
CREATE SERVER pglog FOREIGN DATA WRAPPER file_fdw;
```

```
CREATE FOREIGN TABLE gaiyin (  
    nombre text,  
    numero integer,  
    letra char(1)  
) SERVER pglog  
OPTIONS ( filename 'c:\temp\file.csv', format 'csv' );
```

```
postgres=# explain select * from gaiyin;  
          QUERY PLAN
```

```
-----  
Foreign Scan on gaiyin (cost=0.00..1.10 rows=1 width=44)  
  Foreign File: c:\temp\file.csv  
  Foreign File Size: 36  
(3 filas)
```

<http://wiki.postgresql.org/wiki/SQL/MED>



PALOMINO^{DB}

Proven Database Excellence

Foreign Data Wrapper - mysql_fdw

```
apt-get install libmysqld-dev
```

```
* Compilar como un contrib (make; make install)
```

```
CREATE EXTENSION mysql_fdw;
```

```
CREATE SERVER mysql_srv  
FOREIGN DATA WRAPPER mysql_fdw  
OPTIONS (address 'localhost', port '3306');
```

```
CREATE USER MAPPING  
FOR PUBLIC SERVER mysql_srv  
OPTIONS (username 'root', password 'postgres');
```

```
CREATE FOREIGN TABLE prueba_my  
(texto text, numero int, floaty float)  
SERVER mysql_srv OPTIONS (table 'palomino.prueba');
```

```
CREATE FOREIGN TABLE prueba_my (  
id integer,  
name text,  
address text)  
SERVER mysql_srv  
OPTIONS (table 'palomino.prueba');
```

Fuente: <http://pgsnake.blogspot.com/2011/08/postgresql-91-meet-mysql.html>



PALOMINO^{DB}

Proven Database Excellence

Foreign Data Wrappers (other plugins)

Oracle
Mysql
OBDC
Redis
CouchDB
Filetext
Twitter
LDAP
...

http://wiki.postgresql.org/wiki/Foreign_data_wrappers
<http://repositorio.siu.edu.ar/trac/postgresql/wiki/Dblink>



PALOMINO^{DB}

Proven Database Excellence

Per column collation

```
CREATE TABLE prueba_texto (  
  _id serial PRIMARY KEY,  
  texto_en text COLLATE 'en_US',  
  texto_es text COLLATE 'es_ES'  
);
```

```
SELECT  
  texto_es < ('rendevouz' COLLATE "es_ES")  
FROM prueba_texto;
```

```
SELECT  
  texto_es COLLATE "en_EN" < texto_en  
FROM prueba_texto;
```

```
CREATE INDEX ON prueba  
  (texto COLLATE "en_US");
```

- < 8.4 the collation was strongly configured with the SO, being less flexible.
- In this version we can declare type columns and specify a different collation for each one.
- You can create indexes with a specific collation
 - Be careful, if you use another collation in the query, the index will not be used.



SSI

- Serializable Snapshot Isolation
- Postgres supports now, 3 transactional levels:
 - READ COMMITTED
 - REPEATABLE READ (old SERIALIZABLE)
 - Called SNAPSHOT ISOLATION
 - SERIALIZABLE
- No more “select for update”
- Simplify your code
- It has cost, use predictive locking and could increase the number of rollbacks due conflicts between transactions.

<http://wiki.postgresql.org/wiki/SSI>

<http://wiki.postgresql.org/images/4/4f/SSI-PGConfEU2011.pdf>



PALOMINO^{DB}

Proven Database Excellence

Unlogged tables

- These tables cannot be restored from the WAL in case of crash.
- More performance due of less impact of writing to disk.
- No replication for this tables.
- Truncated in case of crash.
- CLOG contention article: <http://rhaas.blogspot.com/2011/11/hint-bits.html>



Writable Common Table Expressions

- WCTE
- CTE implemented since 8.4
- Advanced info: <http://wiki.postgresql.org/images/d/da/Ciolfi-debug-wcte-2011.pdf>

```
WITH table_temp(i,j,k) AS  
(  
    DELETE FROM prueba WHERE i =9 RETURNING *  
)  
INSERT INTO prueba_del  
SELECT i,j,k  
FROM table_temp;
```

```
WITH RECURSIVE t(a,b) AS (  
    VALUES(0,1)  
    UNION ALL  
    SELECT greatest(a,b), a + b AS a FROM t  
    WHERE b < $1  
)  
SELECT a FROM t;
```



KNN in GiST and GIN

- Nearest values using indexes
 - Increases drastically the performance for Geographical and complex data.
 - pg_trgm can be used for improve LIKE and ILIKE searches
- Operator: “<->”

Example article:

<http://www.palominodb.com/blog/2011/10/12/indexing-text-columns-gist-or-gin-optimize-ilike-using-pgtrgm-postgres-91-part-1>

<http://wiki.postgresql.org/images/4/46/Knn.pdf>



PALOMINO^{DB}

Proven Database Excellence

SE-LINUX support

- Integrated with SE-LINUX.
- The SO permissions are more stronger than database ones.
- Must be enabled explicitly .
- Is based in the SECURITY LABELS creation.
- NOTE: It requires SELINUX libs newer than 2.0.93.

http://wiki.postgresql.org/wiki/SEPostgreSQL_Documentation



PALOMINO^{DB}

Proven Database Excellence

Minor features (ALTER ENUM)

- `ALTER TYPE tipo_enum ADD VALUE new_value [{BEFORE|AFTER} existing_value];`

`CREATE TYPE lista AS ENUM ('tipoa','tipob','tipoc');`

`ALTER TYPE lista ADD VALUE 'tipoa2' AFTER 'tipoa';`



PGXN

<https://github.com/pgxn/pgxn-client>

<http://pgxn.org/faq/>



PALOMINO^{DB}

Proven Database Excellence

General and compatibility changes

- Composite types casts were disabled
 - *valor_compuesto.text*, *text(valor_compuesto)*
- CREATE EXTENSION
- Auto tuning of *wal_buffers*
- GiST and GIN improvements
- INSTEAD OF for triggers (for updateables views setup)
- Windows 64 compilation support
- *standard_conforming_strings* now is 'on' by default
 - Before: 'O\'Connor' , now you should use 'O"Connor' or E'O\'Connor'
- PL/SQL changes in RAISE without parameters
- Inherited tables improvement (Merge Append node – Order of the inherited tables inside the append node instead outside).
- FULL OUTER JOIN could use now Hash nodes.
- Detailed info: <http://www.postgresql.org/docs/9.1/static/release-9-1.html>



Changes (II)

- Replication grants (no more superusers for replictaion)
- Improvement in the security to avoid internal attacks in the replication communication.
- Functional dependencies now are recognized by the GROUP BY (only PK)
 - Before: ERROR: column "tabla.col_pk" must appear in the GROUP BY clause or be used in an aggregate function.



Miscellaneous info

- PLV8 (JavaScript V8 procedural language)
- Postgres-XC <http://postgres-xc.sourceforge.net/>
- PgPool II 3.1
- New versions londiste, plproxy, bucardo and others



PostgreSQL 9.2

Features incorporated:

- * Index only scans
- * Slave replication cascade
- * Oracle style partitioning
- * Materialized Views
- * ...

http://wiki.postgresql.org/wiki/PostgreSQL_9.2_Development_Plan




PALOMINO^{DB}

Proven Database Excellence

Gracias por asistir!!

Presentación hecha en Libre/Open Office





GO...
Save yourself...
I'll keep them...

kill your time on 9GAG.COM



PALOMINO^{DB}

Proven Database Excellence